

Week 15 - Wednesday

CS222

Last time

- What did we talk about last time?
- Review up to Exam 1

Questions?

Project 6

Programming Practice

- Write a function that converts an `int` to a string representation in hexadecimal
- Write a function that converts a string representation of an integer in hexadecimal to an `int`

Review up to Exam 2

Pointers

- A **pointer** is a variable that holds an address
- Often this address is to another variable
- Sometimes it's to a piece of memory that is mapped to file I/O or something else
- Important operations:
 - Reference (**&**) gets the address of something
 - Dereference (*****) gets the contents of a pointer

Declaration of a pointer

- We typically want a pointer that points to a certain kind of thing
- To declare a pointer to a particular type

```
type * name;
```

- Example of a pointer with type **int**:

```
int * pointer;
```

Reference operator

- A fundamental operation is to find the address of a variable
- This is done with the reference operator (`&`)

```
int value = 5;  
int* pointer;  
pointer = &value;  
//pointer has value's address
```

- We usually can't predict what the address of something will be

Dereference operator

- The reference operator doesn't let you do much
- You can get an address, but so what?
- Using the dereference operator, you can read and write the contents of the address

```
int value = 5;  
int* pointer;  
pointer = &value;  
printf("%d", *pointer); //prints 5  
*pointer = 900; //value just changed!
```

Pointer arithmetic

- One of the most powerful (and most dangerous) qualities of pointers in C is that you can take arbitrary offsets in memory
- When you add to (or subtract from) a pointers, it jumps the number of bytes in memory of the size of the type it points to

```
int a = 10;
int b = 20;
int c = 30;
int* value = &b;
value++;
printf("%d", *value); //what does it print?
```

Arrays are pointers too

- An array **is** a pointer
 - It is pre-allocated a fixed amount of memory to point to
 - You can't make it point at something else
- For this reason, you can assign an array directly to a pointer

```
int numbers[] = {3, 5, 7, 11, 13};  
int* value;  
  
value = numbers;  
value = &numbers[0]; //exactly equivalent  
  
//What about the following?  
value = &numbers;
```

Surprisingly, pointers are arrays too

- Well, no, they aren't
- But you can use array subscript notation ([]) to read and write the contents of offsets from an initial pointer

```
int numbers[] = {3, 5, 7, 11, 13};  
int* value = numbers;  
  
printf("%d", value[3]); //prints 11  
printf("%d", *(value + 3)); //prints 11  
value[4] = 19; //changes 13 to 19
```

void pointers

- What if you don't know what you're going to point at?
- You can use a **void***, which is an address to....something!
- You have to cast it to another kind of pointer to use it
- You can't do pointer arithmetic on it
- It's not useful very often

```
char s[] = "Hello World!";
void* address = s;
int* thingy = (int*)address;
printf("%d\n", *thingy);
```

Functions that can change arguments

- In general, data is passed **by value**
- This means that a variable cannot be changed for the function that calls it
- Usually, that's good, since we don't have to worry about functions screwing up our data
- It's annoying if we need a function to return more than one thing, though
- Passing a pointer is equivalent to passing the original data **by reference**

Pointers to pointers

- Just as we can declare a pointer that points at a particular data type, we can declare a pointer to a pointer
- Simply add another star

```
int value = 5;  
int* pointer;  
int** amazingPointer;  
pointer = &value;  
amazingPointer = &pointer;
```

Change `main()` to get command line arguments

- To get the command line values, use the following definition for `main()`

```
int main(int argc, char** argv)
{
    return 0;
}
```

- Is that even allowed?
 - Yes.
- You can name the parameters whatever you want, but `argc` and `argv` are traditional
 - `argc` is the number of arguments (argument count)
 - `argv` are the actual arguments (argument values) as strings

scanf ()

- So far, we have only talked about using **getchar ()** (and command line arguments) for input
- As some of you have discovered, there is a function that parallels **printf ()** called **scanf ()**
- **scanf ()** can read strings, **int** values, **double** values, characters, and anything else you can specify with a % formatting string

```
int number;  
scanf ("%d", &number);
```

Format specifiers

- These are mostly what you would expect, from your experience with `printf()`

Specifier	Type
<code>%d</code>	<code>int</code>
<code>%u</code>	<code>unsigned int</code>
<code>%o</code> <code>%x</code>	<code>unsigned int</code> (in octal for <code>o</code> or hex for <code>x</code>)
<code>%hd</code>	<code>short</code>
<code>%c</code>	<code>char</code>
<code>%s</code>	null-terminated string
<code>%f</code>	<code>float</code>
<code>%lf</code>	<code>double</code>
<code>%Lf</code>	<code>long double</code>

Dynamic Memory Allocation

malloc ()

- Memory can be allocated dynamically using a function called **malloc ()**
 - Similar to using **new** in Java or C++
 - **#include <stdlib.h>** to use **malloc ()**
- Dynamically allocated memory is on the heap
 - It doesn't disappear when a function returns
- To allocate memory, call **malloc ()** with the number of bytes you want
- It returns a pointer to that memory, which you cast to the appropriate type

```
int* data = (int*)malloc(sizeof(int));
```

Allocating arrays

- It is common to allocate an array of values dynamically
- The syntax is exactly the same as allocating a single value, but you multiply the size of the type by the number of elements you want

```
int i = 0;  
int* array = (int*)malloc(sizeof(int)*100);  
for( i = 0; i < 100; i++ )  
    array[i] = i + 1;
```

free()

- C is not garbage collected like Java
- If you allocate something on the stack, it disappears when the function returns
- If you allocate something on the heap, you have to deallocate it with **free()**
- **free()** does not set the pointer to be **NULL**
 - But you can afterwards

```
char* things = (char*)malloc(100) ;  
free(things) ;
```

Ragged Approach

- One way to dynamically allocate a 2D array is to allocate each row individually

```
int** table = (int**)malloc(sizeof(int*)*rows);
int i = 0;

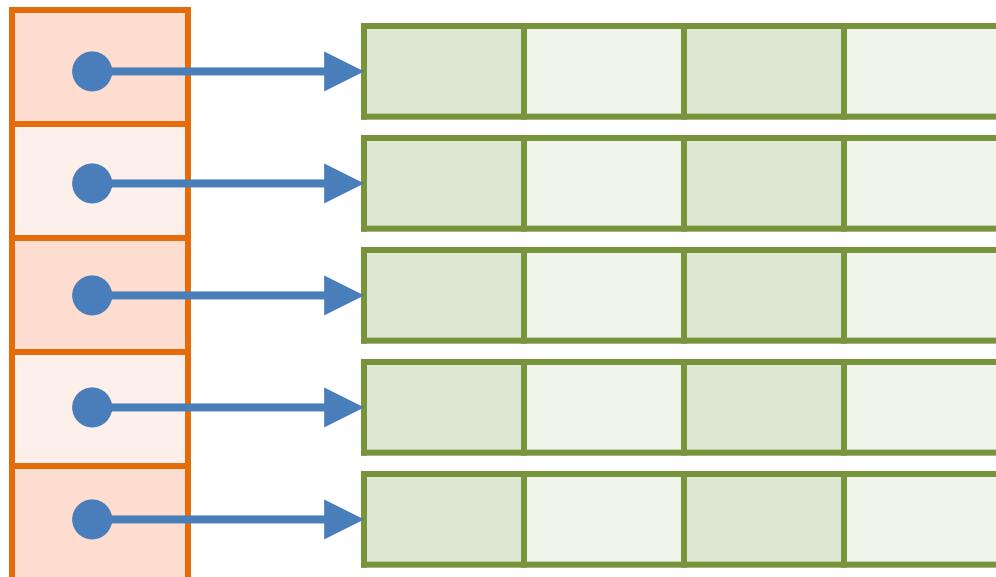
for( i = 0; i < rows; i++ )
    table[i] = (int*)malloc(sizeof(int)*columns);
```

- When finished, you can access **table** like any 2D array

```
table[3][7] = 14;
```

Ragged Approach in memory

table



Chunks of data
that could be
anywhere in
memory

Freeing the Ragged Approach

- To free a 2D array allocated with the Ragged Approach
 - Free each row separately
 - Finally, free the array of rows

```
for( i = 0; i < rows; i++ )  
    free( table[i] );  
  
free( table );
```

Contiguous Approach

- Alternatively, you can allocate the memory for all rows at once
- Then you make each row point to the right place

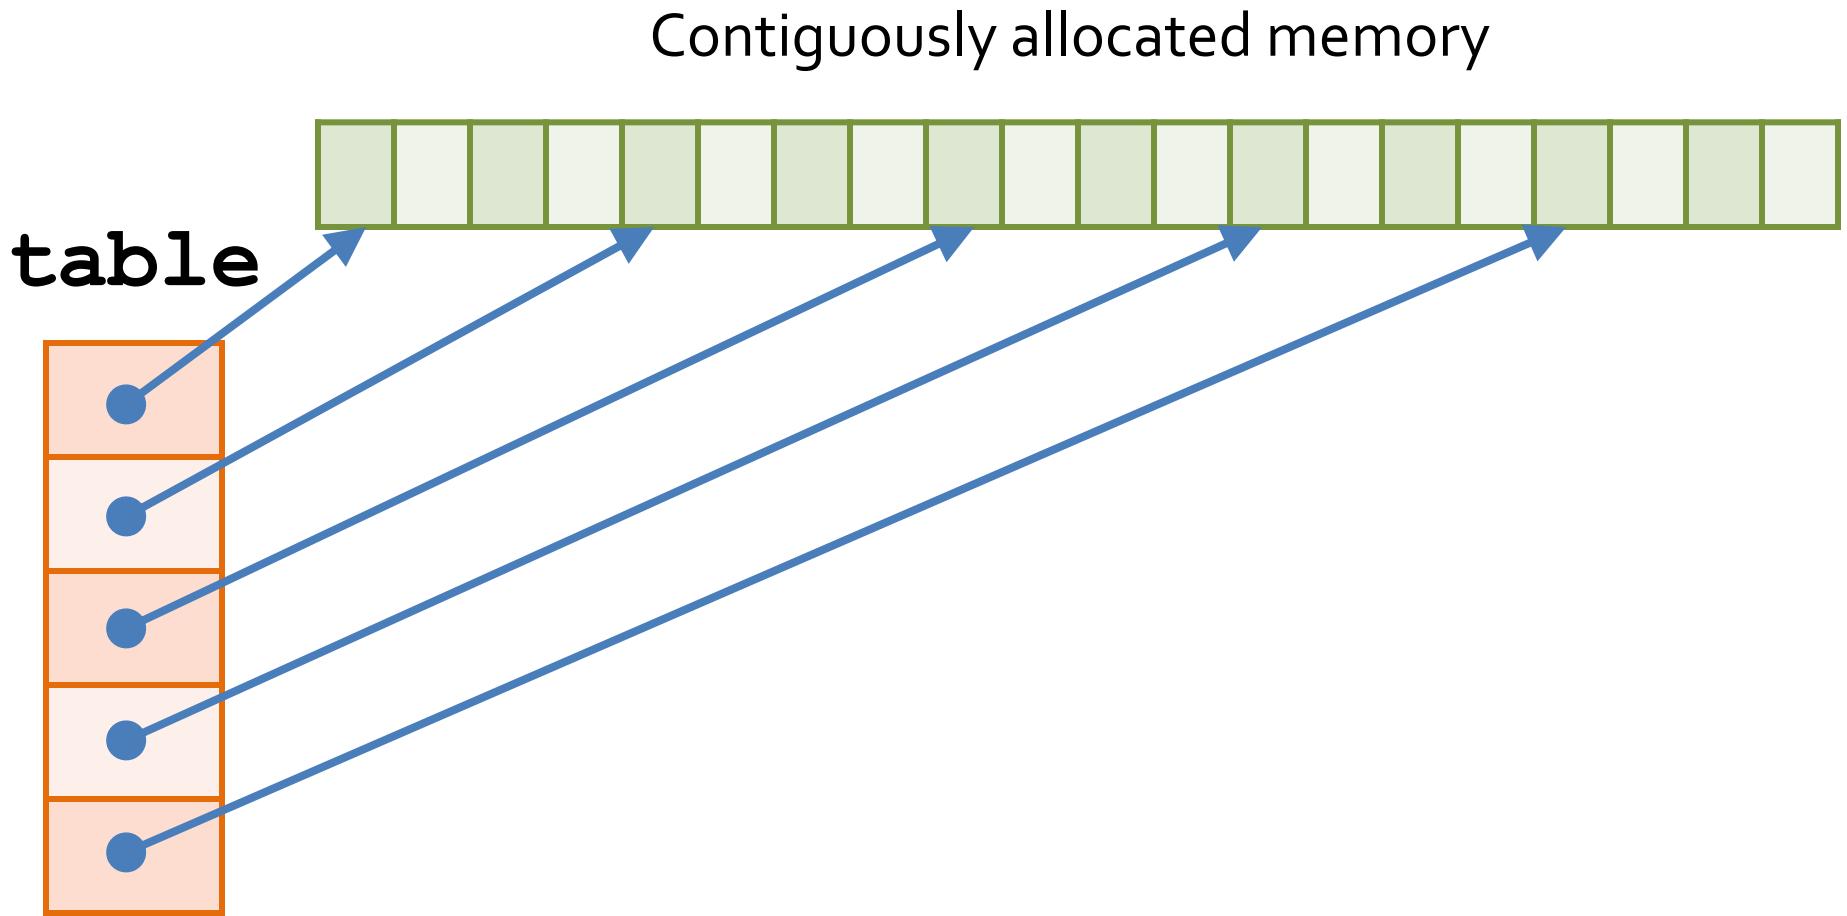
```
int** table = (int**)malloc(sizeof(int*)*rows);
int* data = (int*)malloc(sizeof(int)*rows*columns);
int i = 0;

for( i = 0; i < rows; i++ )
    table[i] = &data[i*columns];
```

- When finished, you can still access **table** like any 2D array

```
table[3][7] = 14;
```

Contiguous Approach in memory



Freeing the Contiguous Approach

- To free a 2D array allocated with the Ragged Approach
 - Free the big block of memory
 - Free the array of rows
 - No loop needed

```
free( table[0] ) ;  
free( table ) ;
```

Comparing the approaches

RAGGED

- Pros
 - Each row can be allocated and freed independently
 - Rows can be shuffled in order with only pointer changes
 - Rows can be different lengths
- Cons
 - Fragmented memory
 - Less locality of reference
 - Requires a loop to free

CONTIGUOUS

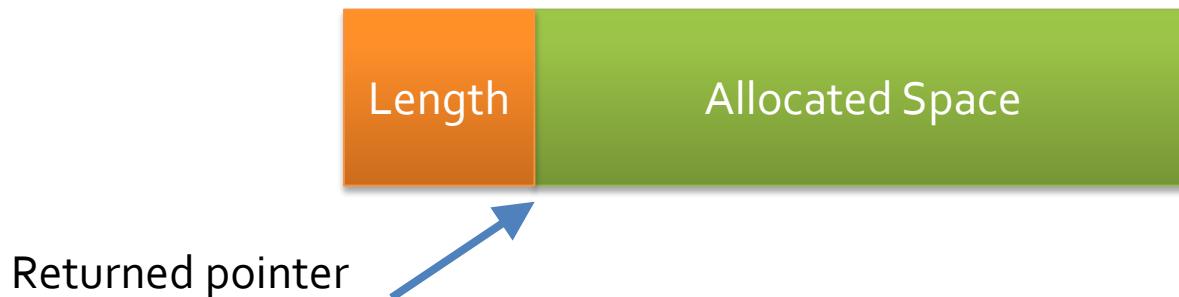
- Pros
 - Better locality of reference
 - Can free the entire thing with two `free()` calls
 - Shuffling rows with pointers is possible, but you also have to keep track of the beginning
- Cons
 - Large allocations are more likely to fail (out of memory)
 - Can't free individual rows

Rules for random numbers

- Include the following headers:
 - `stdlib.h`
 - `time.h`
- Use `rand() % n` to get values between 0 and `n - 1`
- Always call `srand(time(NULL))` before your first call to `rand()`
- Only call `srand()` once per program
 - Seeding multiple times makes no sense and usually makes your output much **less** random

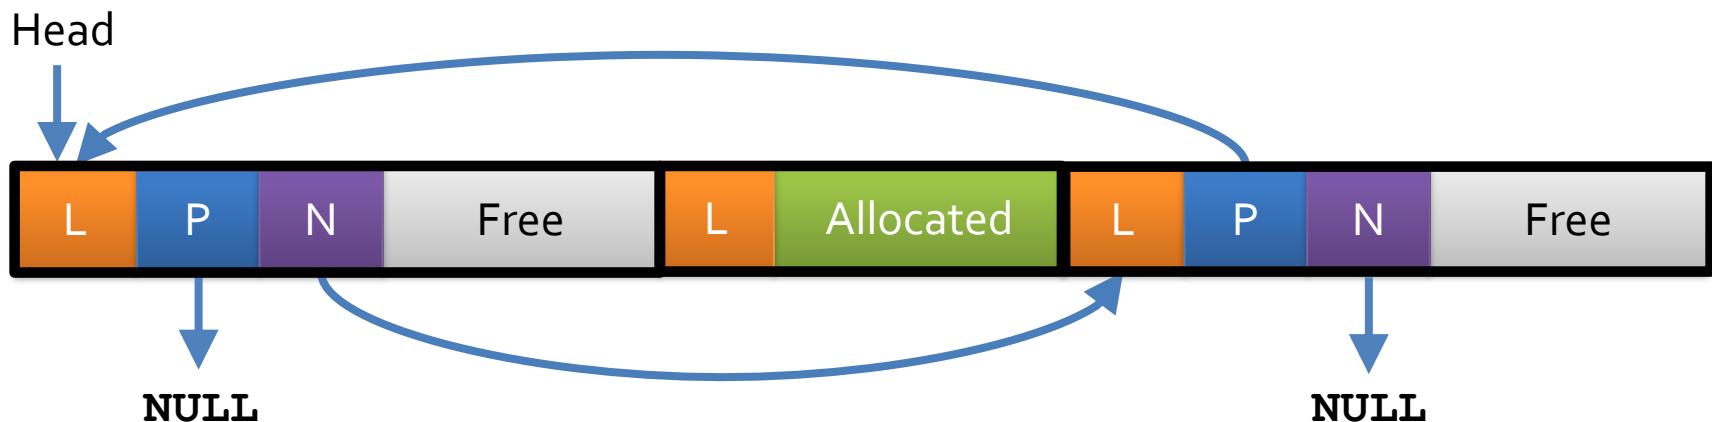
How does malloc() work?

- **malloc()** sees a huge range of free memory when the program starts
- It uses a doubly linked list to keep track of the blocks of free memory, which is perhaps one giant block to begin with
- As you allocate memory, a free block is often split up to make the block you need
- The returned block knows its length
 - The length is usually kept **before** the data that you use



Free list

- Here's a visualization of the free list
- When an item is freed, most implementations will try to coalesce two neighboring free blocks to reduce fragmentation
 - Calling **free()** can be time consuming



Other memory functions

- **void* calloc(size_t items, size_t size);**
 - Clear and allocate **items** items with size **size**
 - Memory is zeroed out
- **void* realloc(void* pointer, size_t size);**
 - Resize a block of memory pointed at by pointer, usually to be larger
 - If there is enough free space at the end, **realloc()** will tack that on
 - Otherwise, it allocates new memory and copies over the old
- **void* alloca(size_t size);**
 - Dynamically allocate memory on the stack (at the end of the current frame)
 - Automatically freed when the function returns
 - You need to **#include <alloca.h>**

Unit testing

- **Unit tests** are tests for a single piece of code in isolation for the rest of your program
- A unit test could test
 - A single method or function
 - A whole class

Test case selection

- Positive test cases
 - Cases that should work, and we can easily check the answer
- Boundary condition test cases
 - Special cases, like deleting the first or last item in a list
 - Sometimes called "corner cases"
- Negative test cases
 - Cases that we know should fail

Regression testing

- When you find a bug, keep the test case that demonstrated the bug
- Include it as part of your test suite
- Always run all your tests when you make a change to your code
- Sometimes a "fix" can break something that was working
 - Especially if you didn't fix the bug correctly the first time

Black box testing

- One philosophy of testing is making **black box tests**
- A black box test takes some input **A** and knows that the output is supposed to be **B**
- It assumes nothing about the internals of the program
- To write black box tests, you come up with a set of input you think covers lots of cases and you run it and see if it works
- In the real world, black box testing can easily be done by a team that did not work on the original development

White box testing

- **White box testing** is the opposite of black box testing
 - Sometimes white box testing is called "clear box testing"
- In white box testing, you can use your knowledge of how the code works to generate tests
- Are there lots of if statements?
 - Write tests that go through all possible branches
- There are white box testing tools that can help you generate tests to exercise all branches
- Which is better, white box or black box testing?

String to integer

- In C, the standard way to convert a string to an `int` is the `atoi()` function
 - `#include <stdlib.h>` to use it

```
#include <stdlib.h>
#include <stdio.h>

int main()
{
    char* value = "3047";
    int x = atoi(value);
    printf("%d\n", x);
    return 0;
}
```

Integer to string

- Oddly enough, this is a stranger situation
- Many systems have a non-standard function (also in **stdlib.h**) called **itoa()**
 - It takes the **int**, a buffer to hold the resulting string, and the base

```
char value[10]; //has to be big enough
int x = 3047;
itoa( x, value, 10 );
```

- The portable way to do this is to use **sprintf()**
 - It's like **printf()** except that it prints things to a string buffer instead of the screen

```
char value[10]; //has to be big enough
int x = 3047;
sprintf( value, "%d", x );
```

Users

- Recall that each user on a Linux system has a unique login name and a unique numerical identifier (the UID)
- Users can belong to one or more groups as well
- Where is this information stored?

Password file

- The system has a password file stored in **/etc/passwd**
- Each line of this file corresponds to one user in the system and has seven fields separated by colons:
 - Login name
 - Encrypted password
 - UID
 - GID (group ID of the first group that the user is a member of)
 - Comment
 - Home directory (where you are when you log in)
 - Login shell (which shell you running when you log in)
- Example:

```
wittmanb:x:1000:100:Barry Wittman:/home/wittmanb:/bin/bash
```

The Linux and Unix solution

- Instead of storing actual passwords, Linux machines store the hash of the passwords
- When someone logs on, the operating system hashes the password and compares it to the stored version
- No one gets to see your original password
 - Not even **root!**

Groups

- Files are associated with a group as well as a user who is owner
- The groups are listed in the /etc/group file
- Each line of this file corresponds to a group and has four fields separated by colons:
 - Group name
 - Encrypted password
 - Often not used
 - Group ID (GID)
 - User list
 - Comma separated
- Example:

```
users:x:100:  
jambit:x:106:claus,felli,frank,harti,markus,martin,mtk,paul
```

Time

- In the systems programming world, there are two different kinds of time that are useful
 - Real time
 - This is also known as wall-clock time or calendar time
 - It's the human notion of time that we're familiar with
 - Process time
 - Process time is the amount of time your process has spent on the CPU
 - There is often no obvious correlation between process time and real time (except that process time is never more than real time elapsed)

time()

- The `time()` function gives back the seconds since the Unix Epoch
- Its signature is:

```
time_t time(time_t* timePointer);
```

- `time_t` is a signed 32-bit or 64-bit integer
- You can pass in a pointer to a `time_t` variable or save the return value (both have the same result)
- Typically we pass in `NULL` and save the return value
- Include `time.h` to use `time()`

```
time_t seconds = time(NULL);
printf("%d seconds have passed since 1970",
      seconds);
```

gettimeofday()

- The **gettimeofday()** function offers a way to get higher precision timing data
- Its signature is:

```
int gettimeofday(struct timeval *tv, struct  
timezone *tz);
```

- The struct **timeval** has a **tv_secs** member which is the same as the return value from **time()**
- It also has a **tv_usec** member which gives microseconds (millionths of a second)
- The **timezone** pointer **tz** is obsolete and should have **NULL** passed into it
- Include **sys/time.h** (not the same as **time.h**) to use this function

Process time

- For optimization purposes, it can be useful to know how much time a process spends running on the CPU
- This time is often broken down into
 - **User time:** the amount of time your program spends executing its own code
 - **System time:** the amount of time spent in kernel mode executing code for your program (memory allocation, page faults, file opening)

The time command

- You can time a program's complete execution by running it with the time command
 - It will give the real time taken, user time, and system time
- Let's say you've got a program called **timewaster**
 - Run it like this:

```
time ./timewaster
```

- Output might be:

```
real 0m4.84s
user 0m1.030s
sys 0m3.43s
```

Structs

- A struct in C is:
 - A collection of one or more variables
 - Possibly of different types
 - Grouped together for convenient handling.
- They were called records in Pascal
- They have similarities to a class in Java
 - Except all fields are public and there are no methods
- Struct declarations are usually global
 - They are outside of `main()` and often in header files

Anatomy of a struct

struct

name

{

type1

member1;

type2

member2;

type3

member3;

...

} ;

Declaring a struct variable

- Type:
 - **struct**
 - The name of the struct
 - The name of the identifier
- You have to put **struct** first

```
struct student bob;
struct student jameel;
struct point start;
struct point end;
```

Accessing members of a struct

- Once you have a struct variable, you can access its members with dot notation (**variable.member**)
 - Members can be read and written

```
struct student bob;
strcpy(bob.name, "Bob Blobberwob");
bob.GPA = 3.7;
bob.ID = 100008;
printf("Bob's GPA: %f\n", bob.GPA);
```

Initializing structs

- There are no constructors for structs in C
- You can initialize each element manually:

```
struct student julio;  
strcpy(julio.name, "Julio Iglesias");  
julio.GPA = 3.9;  
julio.ID = 100009;
```

- Or you can use braces to initialize the entire struct at once:

```
struct student julio =  
{ "Julio Iglesias", 3.9, 100009 };
```

Assigning structs

- It is possible to assign one struct to another

```
struct student julio;
struct student bob;
strcpy(julio.name, "Julio Iglesias");
julio.GPA = 3.9;
julio.ID = 100009;
bob = julio;
```

- Doing so is equivalent to using `memcpy()` to copy the memory of `julio` into the memory of `bob`
- `bob` is still separate memory: it's not like copying references in Java

Dangers with pointers in structs

- With a pointer in a struct, copying the struct will copy the pointer but will not make a copy of the contents
- Changing one struct could change another

```
struct person
{
    char* firstName;
    char* lastName;
};

struct person bob1;
struct person bob2;
```

```
bob1.firstName = strdup("Bob");
bob1.lastName = strdup("Newhart");
bob2 = bob1;
strcpy(bob2.lastName, "Hope");
printf("Name: %s %s\n", bob1.firstName, bob1.lastName);
//prints Bob Hope
```

Arrow notation

- We could dereference a struct pointer and then use the dot to access a member

```
struct student* studentPointer = (struct  
student*) malloc(sizeof(struct student));  
  
(*studentPointer).ID = 3030;
```

- This is cumbersome and requires parentheses
- Because this is a frequent operation, dereference + dot can be written as an arrow (->)

```
studentPointer->ID = 3030;
```

Passing structs to functions

- If you pass a struct directly to a function, you are passing it by value
 - A copy of its contents is made
- It is common to pass a struct by pointer to avoid copying and so that its members can be changed

```
void flip(struct point* value)
{
    double temp = value->x;
    value->x = value->y;
    value->y = temp;
}
```

Gotchas

- **Always** put a semicolon at the end of a struct declaration
- Don't put constructors or methods inside of a struct
 - C doesn't have them
- Assigning one struct to another copies the memory of one into the other
- Pointers to struct variables are usually passed into functions
 - Both for efficiency and so that you can change the data inside

typedef

- The **typedef** command allows you to make an alias for an existing type
- You type **typedef**, the type you want to alias, and then the new name

```
typedef int SUPER_INT;
```

```
SUPER_INT value = 3; //has type int
```

- Don't overuse **typedef**
- It is useful for types like **time_t** which can have different meanings in different systems

typedef with structs

- The **typedef** command is commonly used with structs
 - Often it is built into the struct declaration process
- It allows the programmer to leave off the stupid **struct** keyword when declaring variables

```
typedef struct _wombat
{
    char name[100];
    double weight;
} wombat;
```

- The type defined is actually **struct _wombat**
- We can refer to that type as **wombat**

```
wombat martin;
```

An example linked list node struct

- We'll use this definition for our node for singly linked lists

```
typedef struct _node
{
    int data;
    struct _node* next;
} node;
```

- Somewhere, we will have the following variable to hold the beginning of the list

```
node* head = NULL;
```

Using enum

- To create named constants with different values, type **enum** and then the names of your constants in braces

```
enum { SUNDAY, MONDAY, TUESDAY, WEDNESDAY,  
THURSDAY, FRIDAY, SATURDAY };
```

- Then in your code, you can use these values (which are stored as integers)

```
int day = FRIDAY;  
if( day == SUNDAY )  
    printf("My 'I don't have to run' day");
```

Specifying values

- You can even specify the values in the `enum`

```
enum { ANIMAL = 7, MINERAL = 9, VEGETABLE = 11 };
```

- If you assign values, it is possible to make two or more of the constants have the same value (usually bad)
- A common reason that values are assigned is so that you can do bitwise combinations of values

```
enum { PEPPERONI = 1, SAUSAGE = 2, BACON = 4,  
MUSHROOMS = 8, PEPPER = 16, ONIONS = 32, OLIVES  
= 64, EXTRA_CHEESE = 128 };
```

```
int toppings = PEPPERONI | ONIONS | MUSHROOMS;
```

An example doubly linked list node struct

- We'll use this definition for our node for doubly linked lists

```
typedef struct _node
{
    int data;
    struct _node* next;
    struct _node* previous;
} node;
```

- Somewhere, we will have the following variables to hold the beginning and ending of the list

```
node* head = NULL;
node* tail = NULL;
```

Bit fields in a struct

- You can define a struct and define how many bits wide each element is
 - It only works for integral types, and it makes the most sense for **unsigned int**
 - Give the number of bits it uses after a colon
 - The bits can't be larger than the size the type would normally have
 - You can have unnamed fields for padding purposes

```
typedef struct _toppings
{
    unsigned pepperoni : 1;
    unsigned sausage   : 1;
    unsigned onions   : 1;
    unsigned peppers  : 1;
    unsigned mushrooms: 1;
    unsigned sauce    : 1;
    unsigned cheese   : 2;
    //goes from no cheese to triple cheese
} toppings;
```

Struct size and padding

- Remember that structs are always padded out to multiples of 4 bytes (well, of `int` size)
 - Unless you use compiler specific statements to change byte packing
- After the last bit field, there will be empty space up to the nearest 4 byte boundary
- You can mix bit field members and non-bit field members in a struct
 - Whenever you switch, it will pad out to 4 bytes
 - You can also have 0 bit fields which also pad out to 4 bytes

Declaring unions

- Unions look like structs
 - Put the keyword **union** in place of **struct**

```
union Congressperson
{
    int district;    //representatives
    char state[15]; //senators
};
```

- There isn't a separate district and a state
 - There's only space for the larger one
 - In this case, 15 bytes (rounded up to 16) is the larger one

Practice

- Write a function that finds the median of an array
 - You'll have to sort it
- Write a function that, given a string, creates a dynamically allocated chunk of memory containing the string reversed
- Write a function that will delete an element from the doubly linked list struct given in earlier slides
- Write a program that counts the **total** number of characters in all the arguments passed in through the command line
 - Ignore **argv[0]**
- Write a program to "encrypt" a file by writing a new file with exactly the same contents, except that each byte in the file is inverted
 - Old byte: **x**
 - New byte: **255 - x**

Quiz

Upcoming

Next time...

- Review after Exam 2
- Lab 15

Reminders

- Study for Final Exam
 - 11:00am - 2:00pm, Tuesday, 5/08/2018 (Section A)
 - 11:00am - 2:00pm, Monday, 5/07/2018 (Section B)
- Finish Project 6
 - Due this Friday before midnight