

Week 8 - Wednesday

CS222

# Last time

---

- What did we talk about last time?
- Software engineering
- Testing
- GDB

# Questions?

# Project 4

# Quotes

*Good design adds value faster than it adds cost.*

Thomas C. Gale

# GDB

- GDB (the GNU Debugger) is a debugger available on Linux and Unix systems
- It is a command line utility, but it still has almost all the power that the Eclipse debugger does:
  - Setting breakpoints
  - Stepping through lines of code
  - Examining the values of variables at run time
- It supports C, C++, Objective-C, Java, and other languages

# Prerequisites

- C doesn't run in a virtual machine
- To use GDB, you have to compile your program in a way that adds special debugging information to the executable
- To do so, add the **-ggdb** flag to your compilation

```
gcc -ggdb program.c -o program
```

- Note: You will not need to do this on Friday's lab

# Starting GDB

- The easiest way to run GDB is to have it start up a program
- Assuming your executable is called **program**, you might do it like this:

```
gdb program
```

- It is also possible to attach GDB to a program that is running already, but you have to know its PID
- You can also run GDB on a program that has died, using the core file (which is why they exist)



# Basic GDB commands

| Command               | Shortcut  | Description   |
|-----------------------|-----------|---|
| <b>run</b>            | <b>r</b>  | Start the program running                           |
| <b>list 135</b>       | <b>l</b>  | List the code near line 135                         |
| <b>list function</b>  | <b>l</b>  | List the code near the start of <b>function()</b>   |
| <b>print variable</b> | <b>p</b>  | Print the value of an expression                    |
| <b>backtrace</b>      | <b>bt</b> | List a stack trace                                  |
| <b>break 29</b>       | <b>b</b>  | Set a breakpoint on line 29                         |
| <b>break function</b> | <b>b</b>  | Set a breakpoint at the start of <b>function()</b>  |
| <b>continue</b>       | <b>c</b>  | Start running again after stopping at a breakpoint  |
| <b>next</b>           | <b>n</b>  | Execute next line of code, skipping over a function |
| <b>step</b>           | <b>s</b>  | Execute next line of code, stepping into a function |
| <b>quit</b>           | <b>q</b>  | Quit using GDB                                      |

# Some String Issues

# A few final string issues

- What if you have a number and want a string version of it?
  - In Java:

```
int x = 3047;  
String value = "" + x;           //quick way  
value = Integer.toString(x);    //fussy way
```

- What if you have a string that gives a numerical representation and you want the number it represents?
  - In Java:

```
String value = "3047";  
int x = Integer.parseInt(value);
```

# String to integer

- In C, the standard way to convert a string to an **int** is the **atoi()** function
  - **#include <stdlib.h>** to use it

```
#include <stdlib.h>
#include <stdio.h>

int main()
{
    char* value = "3047";
    int x = atoi(value);
    printf("%d\n", x);
    return 0;
}
```

# Implementing `atoi()`

- Now it's our turn to implement `atoi()`
  - Signature:

```
int atoi(char* number);
```

# Integer to string

- Oddly enough, this is a stranger situation
- Many systems have a non-standard function (also in `stdlib.h`) called `itoa()`
  - It takes the `int`, a buffer to hold the resulting string, and the base

```
char value[10]; //has to be big enough
int x = 3047;
itoa( x, value, 10 );
```

- The portable way to do this is to use `sprintf()`
  - It's like `printf()` except that it prints things to a string buffer instead of the screen

```
char value[10]; //has to be big enough
int x = 3047;
sprintf( value, "%d", x );
```

# Implementing `itoa()`

- Now it's our turn to implement `itoa()`
- We'll restrict ourselves to a base 10 version
  - Signature:

```
char* itoa(int number, char* string);
```

- It returns the pointer to the string in case this call is nested inside of another call (to `strcat()` or `printf()` or whatever)

# Users and Groups



# Users

- Recall that each user on a Linux system has a unique login name and a unique numerical identifier (the UID)
- Users can belong to one or more groups as well
- Where is this information stored?

# Password file

- The system has a password file stored in **/etc/passwd**
- Each line of this file corresponds to one user in the system and has seven fields separated by colons:
  - Login name
  - Encrypted password
  - UID
  - GID (group ID of the first group that the user is a member of)
  - Comment
  - Home directory (where you are when you log in)
  - Login shell (which shell you running when you log in)
- Example:

```
wittmanb:x:1000:100:Barry Wittman:/home/wittmanb:/bin/bash
```

# Catch-22

- Your computer needs to be able read the password file to check passwords
- But, even **root** shouldn't be able to read everyone's passwords
- Hash functions to the rescue!

# Cryptographic hash functions

- Take a long message and turn it into a short digest
- Different from hash functions used for hash tables
- Lots of interesting properties (lots more than these):

## Avalanching

- A small change in the message should make a big change in the digest

## Preimage Resistance

- Given a digest, should be hard to find a message that would produce it

## Collision Resistance

- Should be hard to find two messages that hash to the same digest (collision)

# The Linux and Unix solution

- Instead of storing actual passwords, Linux machines store the hash of the passwords
- When someone logs on, the operating system hashes the password and compares it to the stored version
- No one gets to see your original password
  - Not even **root**!

# Back to the password file

- Inside the password file, we have encrypted passwords
- Everyone's password is safe after all

| Login Name | Password Hash |
|------------|---------------|
| ahmad      | IfW{6Soo      |
| baili      | 853aE90f      |
| carmen     | D390&063      |
| deepak     | CWc^Q3Ge      |
| erica      | e[6s_N*X1     |

# Shadow password file

- Even though the password is disguised, it is unwise to let it be visible to everyone
  - Given a password digest (the hashed version) and lots of time, it is possible to figure out the password
- It's useful for the password file to be readable by everyone so that all users on a machine are known to all others
- A shadow password file stores the encrypted password and is readable only by privileged users
  - **`/etc/shadow`**

# Changing your password

- Amid all this discussion, it might be useful to know how to change your password
- I **don't** recommend that you do change your password
  - I'm honestly not sure how doing so will interact with your Active Directory (Windows) password
- The command is **passwd**

```
Changing password for wittmanb.  
(current) UNIX password:  
Enter new UNIX password:  
Retype new UNIX password:  
passwd: password updated successfully
```



# Changing the owner of a file

- You recall that we can change permissions for who can read, write, and execute a file using **chmod**
- But **chmod** depends on who the owner is
- What if you want someone else to be the owner of a file?
- The **chown** command can let you do that
- If I want my file **stuff.txt** to be owned by Dr. Leap, I would use the following command

```
chown leap stuff.txt
```

- On most systems, **chown** only works if you are **root**

# Groups

- Files are associated with a group as well as a user who is owner
- The groups are listed in the **/etc/group** file
- Each line of this file corresponds to a group and has four fields separated by colons:
  - Group name
  - Encrypted password
    - Often not used
  - Group ID (GID)
  - User list
    - Comma separated
- Example:

```
users:x:100:  
jambit:x:106:claus,felli,frank,harti,markus,martin,mtk,paul
```

# Creating a group

- If you want to create a group, you have to be **root**
- If you're **root** (or using **sudo**), you can use the **groupadd** command
- To create the **awesome** group as **root**:

```
groupadd awesome
```

- Or using **sudo**:

```
sudo groupadd awesome
```

# Adding a user to a group

- Again, you have to be **root** to add a user to a group
- Use the **useradd** command
- To add user **wittmanb** to the **awesome** group as **root**:

```
useradd -g awesome wittmanb
```

- Or using **sudo**:

```
sudo useradd -g awesome wittmanb
```

# Changing the group for a file

- When you create a file, it is associated with some default group that you belong to
- You can use the **chgrp** command to change to another group that you belong to

```
chgrp awesome file.txt
```

- If you are root, you can use the **chown** command to change the group, using a colon

```
chown :awesome file.txt
```

# Quiz

# Upcoming

# Next time...

---

- Time from the Linux perspective
- Lab 8



# Reminders

---

- Keep working on Project 4
- Read LPI Chapter 10